

# Data Engineer — Step 3

## Learn ETL & Data Warehousing

CareerCompass detailed study notes

### Learning objective

Treat this step as a capability rather than a list of keywords. You should be able to explain the underlying idea, follow the normal workflow, build a small example, handle common failures, and show evidence of the skill. This step is part of the Data Engineer roadmap.

#### 1. Core mental model

For **Learn ETL & Data Warehousing**, think in terms of inputs → decisions/transformations → outputs → validation/feedback. Identify what enters the system, what rules or tools act on it, what result is produced, and what evidence proves the result is correct. This model keeps learning connected to real work.

#### 2. Detailed concept map

##### 1. Purpose and scope of learn etl & data warehousing

Study Purpose and scope of learn etl & data warehousing by first defining the important terms, then tracing one concrete example from start to finish. Rebuild the example without copying. Change one requirement and observe what must change. Finally, deliberately introduce a mistake and diagnose it. For this step, the goal is understanding the relationship between the parts, not memorising isolated commands.

##### 2. terminology and component relationships

Study terminology and component relationships by first defining the important terms, then tracing one concrete example from start to finish. Rebuild the example without copying. Change one requirement and observe what must change. Finally, deliberately introduce a mistake and diagnose it. For this step, the goal is understanding the relationship between the parts, not memorising isolated commands.

##### 3. normal workflow from input to result

Study normal workflow from input to result by first defining the important terms, then tracing one concrete example from start to finish. Rebuild the example without copying. Change one requirement and observe what must change. Finally, deliberately introduce a mistake and diagnose it. For this step, the goal is understanding the relationship between the parts, not memorising isolated commands.

##### 4. edge cases, failure modes and debugging

Study edge cases, failure modes and debugging by first defining the important terms, then tracing one concrete example from start to finish. Rebuild the example without copying. Change one requirement and observe what must change. Finally, deliberately introduce a mistake and diagnose it. For this step, the goal is understanding the relationship between the parts, not memorising isolated commands.

##### 5. practical implementation/design decisions

Study practical implementation/design decisions by first defining the important terms, then tracing one concrete example from start to finish. Rebuild the example without copying. Change one requirement and observe what must change. Finally, deliberately introduce a mistake and diagnose it. For this step, the goal is understanding the relationship between the parts, not memorising isolated commands.

##### 6. quality, maintainability and real-world trade-offs

Study quality, maintainability and real-world trade-offs by first defining the important terms, then tracing one concrete example from start to finish. Rebuild the example without copying. Change one requirement and observe what must change. Finally, deliberately introduce a mistake and diagnose it. For this step, the goal is understanding the relationship between the parts, not memorising isolated commands.

### 3. Practical learning sequence

- A. Learn the vocabulary and the purpose of each component.
- B. Follow one small worked example.
- C. Recreate it from an empty project/file.
- D. Modify one requirement.
- E. Test normal and edge cases.
- F. Debug an intentional failure.
- G. Save the final artifact and document the decisions.

### 4. Worked practice

Build a small practice task directly related to **Learn ETL & Data Warehousing**. Write the requirement first. List inputs, expected outputs and constraints. Choose the simplest suitable approach, implement incrementally, and record at least three test cases. Include one negative/edge case and explain how you diagnosed any failure.

### 5. Mini-project / portfolio evidence

Create a portfolio artifact for this step: a working implementation or design, a README/problem statement, setup or usage instructions, key decisions, three checks/tests, and a short troubleshooting note. The artifact should demonstrate the actual skill in **Learn ETL & Data Warehousing**, not just a copied tutorial.

### 6. Common mistakes

- memorising syntax without understanding the model
- copying tutorials without rebuilding from scratch
- ignoring edge cases and failure states
- choosing tools before understanding the problem
- not reading official documentation when behaviour is unclear
- being unable to explain why a decision was made

### 7. Troubleshooting method

Reproduce → isolate → inspect inputs/state → form one hypothesis → make one change → retest → document the cause and fix. For software use logs, stack traces, browser/network tools and minimal reproductions. For design use task evidence, hierarchy, consistency and accessibility checks rather than personal preference alone.

### 8. Assessment questions

- What problem does learn etl & data warehousing solve?
- What are the three most important concepts here?
- How do you verify that the result is correct?
- What are two common failure modes?
- When would you choose a simpler alternative?
- How would you demonstrate this skill in an interview?

### 9. Mastery checklist

- Explain it without notes.
- Build a small example from scratch.
- Explain major decisions.
- Debug a broken example.
- Handle three edge cases.
- Produce portfolio evidence.
- Answer the assessment questions.

## 10. Recommended documentation

- **Python Documentation** — <https://docs.python.org/3/>
- **pandas Documentation** — <https://pandas.pydata.org/docs/>
- **Khan Academy Statistics** — <https://www.khanacademy.org/statistics-probability>
- **Microsoft Learn Data Analyst** — <https://learn.microsoft.com/en-us/training/career-paths/data-analyst>
- **Power BI Learning** — <https://learn.microsoft.com/en-us/training/paths/data-analytics-microsoft/>
- **Tableau Training** — <https://www.tableau.com/learn/training>

## 11. Direct video learning

- Direct watch URL — [https://www.youtube.com/watch?v=PHsC\\_t0j1dU](https://www.youtube.com/watch?v=PHsC_t0j1dU)

# Deep-Dive Appendix

## A. Mechanism and reasoning

Trace a complete example of learn etl & data warehousing from beginning to end. At every boundary ask: what data/state enters here, what assumptions are being made, what transformation happens, and what observable evidence proves success? Then identify where a failure could occur and what evidence would distinguish a logic problem from a configuration, data, environment or user problem.

## B. Compare alternatives

Compare two legitimate approaches to the same small problem. Record complexity, maintainability, performance/usability, ecosystem/tooling, failure modes and the situations in which each approach is appropriate. The purpose is to learn decision-making rather than memorise a single tool.

## C. Extended practice

Add one realistic constraint to your learn etl & data warehousing practice task: changing requirements, imperfect input, limited time, accessibility, security, scale, deployment or collaboration. Explain how the constraint changed your implementation/design and why.

## D. Revision sheet

Write one page from memory: definitions, workflow, one diagram/flow, one worked example, three mistakes, three checks and three interview questions. Revisit any section you cannot explain clearly.

## E. Final mastery test

Without opening a tutorial, teach learn etl & data warehousing for five minutes, build a small example, intentionally break it, diagnose the failure, and improve the result. Mark the step complete only when you can explain, build, break, diagnose and improve.